

# Deep Learning I — Final Project Report

**Course:** Deep Learning I: Neural Networks (Semester 4) · Section A · Group 1

**Group members (8):** Shreyansh Arora (24BCS10252), Bhavya Patel (24BCS10420), Rudra Longaonkar (24BCS10312), Mohammed Rehan (24BCS10620), Avishkar Chavan (24BCS10065), Milap Kothari (24BCS10004), Soham Ambore (24BCS10258), Shubh Shukla (24BCS10093).

## Abstract

We present two studies. **Task 1 (applied)** forecasts air temperature **12 hours ahead** from the Jena Climate dataset, comparing a moving-average/persistence baseline, a vanilla RNN, and an LSTM, and analysing the effect of the lookback-window length. **Task 2 (experimental)** trains a small VGG-style CNN on CIFAR-10 (82.2% test accuracy) and visualises its first-layer filters and feature maps at an early vs. a late layer to show how representations grow more abstract with depth.

## Task 1 — Time Series Forecasting (Weather)

### 1.1 Problem & data

Predict the temperature ( $T$  (degC)) **12 hours ahead** from the Jena Climate dataset (2009–2016, 10-min readings resampled to hourly → ~70k points; ~49k supervised windows). Strong daily seasonality is visible in the one-week zoom (Fig. 1). Data is split **chronologically** 70/15/15 and standardised using **train** statistics only (no leakage). Inputs are sliding windows of `LOOKBACK` hours → the value 12 h after the window.




Fig. 1 — Full hourly series (2009–2016) and a one-week zoom showing the daily cycle.

### 1.2 Methods

- **Baseline** — persistence (forecast = last observed value) and moving average of the window. No training.
- **Vanilla RNN** — `nn.RNN`, hidden=64, last hidden state → linear head.
- **LSTM** — `nn.LSTM`, hidden=64, same head. Identical training (Adam, MSE, 12 epochs) so the only difference is the recurrent cell. All metrics (MAE, RMSE) are reported in °C after de-normalisation.

### 1.3 Results

| Model          | MAE (°C)    | RMSE (°C)   |
|----------------|-------------|-------------|
| <b>LSTM</b>    | <b>2.19</b> | <b>2.81</b> |
| Vanilla RNN    | 2.22        | 2.86        |
| Moving average | 2.96        | 3.73        |
| Persistence    | 4.18        | 5.41        |

The **LSTM is best**, cutting RMSE by ~48% versus the persistence baseline (2.81 vs 5.41). At a 12-hour horizon persistence is badly out of phase — it predicts the temperature from 12 h earlier, i.e. roughly the opposite point of the day/night cycle (Fig. 3, light blue), which is exactly why the recurrent models win so clearly here.




Fig. 2 — Training/validation MSE (normalised). Both converge; the LSTM reaches a lower val loss.




Fig. 3 — 12-h-ahead forecast over a one-week test slice. RNN/LSTM track the daily cycle; persistence is a full half-day out of phase.

### 1.4 Ablation — lookback window length

We retrain RNN and LSTM for lookbacks [6, 12, 24, 48, 96] h (Fig. 4).

| Lookback (h) | RNN RMSE | LSTM RMSE |
|--------------|----------|-----------|
| 6            | 3.92     | 4.16      |
| 12           | 3.06     | 2.99      |
| 24           | 2.96     | 2.86      |

| Lookback (h) | RNN RMSE | LSTM RMSE   |
|--------------|----------|-------------|
| 48           | 2.88     | 2.82        |
| 96           | 2.88     | <b>2.74</b> |




Fig. 4 — Test RMSE vs lookback window length.

**Analysis.** A 6-hour window is **too short** — it can't see the daily cycle, and error is worst ( $\approx 3.9\text{--}4.2$  °C). Error drops sharply as the window approaches  $\sim 24$  h (one full day), then **flattens**. Beyond that the two cells diverge: the **LSTM keeps improving** out to 96 h ( $2.82 \rightarrow 2.74$ ) because its gates retain useful long-range context, whereas the **vanilla RNN plateaus and slightly worsens** ( $2.875 \rightarrow 2.881$ ), consistent with its weaker long-sequence memory. Practical takeaway:  $\sim 24\text{--}48$  h is the sweet spot for the RNN; the LSTM benefits from longer windows.

## Task 2 — CNN Filter Visualization (CIFAR-10)

### 2.1 Problem & model

Train a 4-conv VGG-style CNN (conv1 32  $\rightarrow$  conv2 64  $\rightarrow$  conv3 128  $\rightarrow$  conv4 128, two max-pools, FC head; **2,340,810 params**) on CIFAR-10 with light augmentation (random crop + flip) and channel normalisation.




Fig. 5 — Sample CIFAR-10 images.

### 2.2 Results

**Final test accuracy: 82.2%** (12 epochs). Curves are healthy with a small train/test gap.




Fig. 6 — Training/test loss and accuracy.




Fig. 7 — Test-set confusion matrix.

Per-class accuracy is highest for visually distinct classes (ship 92%, plane 90%, truck 89%, deer 87%) and lowest for visually similar animals (cat 61%, bird 70%, dog 75%). Most errors fall between semantically close classes — notably **cat**↔**dog** (128 + 92 misclassified) and **car**↔**truck** (77) — which is the expected CIFAR-10 failure mode.

## 2.3 Filter & feature-map visualization




Fig. 8 — The 32 first-layer (conv1) filters: low-level edge / colour-opponent detectors.




Fig. 9 — Feature maps for one input. Early (conv1, 32×32) are sharp and retain the object outline; late (conv4, 16×16) are coarse, sparse, and abstract.

## 2.4 Analysis — abstraction with depth

The network builds a hierarchy: pixels → edges/colours (conv1) → textures/parts → abstract, class-discriminative features (conv4). Early maps answer *where* (spatial detail); deep maps answer *what* (semantics), trading resolution for meaning. Many conv4 channels are silent for a given image (sparse, selective activations) — a hallmark of learned high-level detectors. This hierarchy is *why* deep stacks of small convolutions classify well, and it is consistent with the confusion-matrix errors clustering among visually similar classes.

## Conclusion

For 12-h-ahead temperature forecasting, **LSTM > RNN » moving average » persistence**, and the lookback sweep shows error falling until ~one daily period then flattening, with the LSTM uniquely benefiting from longer context. The CNN study makes the depth→abstraction story concrete and visual at 82.2% accuracy.

**Limitations / future work:** multi-step (sequence) forecasting and exogenous features (pressure, humidity) for Task 1; deeper nets / Grad-CAM for Task 2.

## Member contributions

| Member           | Roll No    | Email                               | Contribution                                               |
|------------------|------------|-------------------------------------|------------------------------------------------------------|
| Shreyansh Arora  | 24BCS10252 | shreyansh.24bcs10252@sst.scaler.com | Project coordination; Task 1 data pipeline & preprocessing |
| Bhavya Patel     | 24BCS10420 | bhavya.24bcs10420@sst.scaler.com    | Task 1 RNN/LSTM model implementation & training            |
| Rudra Longaonkar | 24BCS10312 | rudra.24bcs10312@sst.scaler.com     | Task 1 lookback-window ablation & forecast analysis        |
| Mohammed Rehan   | 24BCS10620 | mohammed.24bcs10620@sst.scaler.com  | Task 1 baselines, metrics (MAE/ RMSE) & evaluation         |
| Avishkar Chavan  | 24BCS10065 | avishkar.24bcs10065@sst.scaler.com  | Task 2 CNN architecture & training                         |
| Milap Kothari    | 24BCS10004 | milap.24bcs10004@sst.scaler.com     | Task 2 evaluation: confusion matrix & training curves      |
| Soham Ambore     | 24BCS10258 | soham.24bcs10258@sst.scaler.com     | Task 2 filter & feature-map visualization                  |
| Shubh Shukla     | 24BCS10093 | shubh.24bcs10093@sst.scaler.com     | Report writing & slide deck preparation                    |

*Note: every member ran both notebooks end-to-end and can explain any part for the viva.*

## References

- Jena Climate dataset (Max Planck Institute for Biogeochemistry).
- Krizhevsky, *Learning Multiple Layers of Features from Tiny Images* (CIFAR-10), 2009.
- Hochreiter & Schmidhuber, *Long Short-Term Memory*, Neural Computation, 1997.